

# Making Kruskal's algorithm fast

## COMS20017 (Algorithms and Data)

John Lapinskas, University of Bristol

# Implementing Kruskal's algorithm

---

**Algorithm:** KRUSKAL

---

**Input** : Connected weighted graph  $G = ((V, E), w)$  in adjacency list form.

**Output** : A minimum spanning tree for  $G$ .

- 1 Sort the edges by weight as  $e_1, \dots, e_m$ , with  $w(e_1) \leq \dots \leq w(e_m)$ .
  - 2 Let  $T \leftarrow (V, \emptyset)$  be the empty tree on  $V$ .
  - 3 **for**  $i = 1$  to  $m$  **do**
  - 4     **if**  $T + e_i$  has no cycles **then**
  - 5         Let  $T \leftarrow T + e_i$ .
  - 6 Return  $T$ .
- 

Lines 1, 2 and 6 take  $O(|E| \log |E|)$  time, and lines 3–5 repeat  $|E|$  times.

We *could* implement line 4 with BFS... but this would take  $\Theta(|E|)$  time, giving us a worst-case running time of  $\Theta(|E|^2)$ . That's bad.

# Implementing Kruskal's algorithm: Take 2

**Idea:** Joining two tree components with an edge will never add a cycle, and adding an edge inside a tree component will always add one.

So when we consider an edge  $e_i$  to  $T$ , we just need to make sure both endpoints aren't in the same component — this implementation will work:

---

## Algorithm: KRUSKAL

---

**Input** : Connected weighted graph  $G = ((V, E), w)$  in adjacency list form.

**Output** : A minimum spanning tree for  $G$ .

- 1 Sort the edges by weight as  $e_1, \dots, e_m$ , with  $w(e_1) \leq \dots \leq w(e_m)$ .
  - 2 Let  $T \leftarrow (V, \emptyset)$  be the empty tree on  $V$ .
  - 3 Let  $\mathcal{C} \leftarrow$  the set of  $T$ 's components.
  - 4 **for**  $i = 1$  to  $m$  **do**
  - 5     Let  $C_1$  and  $C_2$  be the components containing  $e_i$ 's endpoints in  $\mathcal{C}$ .
  - 6     **if**  $C_1 \neq C_2$  **then**
  - 7         Let  $T \leftarrow T + e_i$ .
  - 8         Merge  $C_1$  and  $C_2$  in  $\mathcal{C}$ .
  - 9 Return  $T$ .
-

# The key problem

---

## Algorithm: KRUSKAL

---

**Input** : Connected weighted graph  $G = ((V, E), w)$  in adjacency list form.

**Output** : A minimum spanning tree for  $G$ .

- 1 Sort the edges by weight as  $e_1, \dots, e_m$ , with  $w(e_1) \leq \dots \leq w(e_m)$ .
  - 2 Let  $T \leftarrow (V, \emptyset)$  be the empty tree on  $V$ .
  - 3 Let  $\mathcal{C} \leftarrow$  the set of  $T$ 's components.
  - 4 **for**  $i = 1$  to  $m$  **do**
  - 5     Let  $C_1$  and  $C_2$  be the components containing  $e_i$ 's endpoints in  $\mathcal{C}$ .
  - 6     **if**  $C_1 \neq C_2$  **then**
  - 7         Let  $T \leftarrow T + e_i$ .
  - 8         Merge  $C_1$  and  $C_2$  in  $\mathcal{C}$ .
  - 9 Return  $T$ .
- 

But how do we implement  $\mathcal{C}$ ?

A linked list for each component? Then merging will take  $O(1)$  time, but finding  $C_1$  and  $C_2$  could take  $\Omega(|V|)$  time, giving a runtime of  $\Omega(|V||E|)$ .

An array for each component? Then finding  $C_1$  and  $C_2$  will take  $O(1)$  time, but merging will take  $\Omega(|V|)$ , so we still get  $\Omega(|V||E|)$  overall...

# The solution

We need to use a **union-find** data structure, also known as a **disjoint-set** or **merge-find** data structure. It supports the following operations:

- $\text{MakeUnionFind}(X)$ : Makes a new union-find data structure containing a 1-element set  $\{x\}$  for each element  $x \in X$ .
- $\text{Union}(x, y)$ : Merge the set containing  $x$  and the set containing  $y$ .
- $\text{FindSet}(x)$ : Returns a unique identifier for the set containing  $x$ .

1	2	3	4	5	6
$\{v_1\}$	$\{v_2\}$	$\{v_3\}$	$\{v_4\}$	$\{v_5\}$	$\{v_6\}$

$\text{MakeUnionFind}(v_1, v_2, v_3, v_4, v_5, v_6);$

# The solution

We need to use a **union-find** data structure, also known as a **disjoint-set** or **merge-find** data structure. It supports the following operations:

- $\text{MakeUnionFind}(X)$ : Makes a new union-find data structure containing a 1-element set  $\{x\}$  for each element  $x \in X$ .
- $\text{Union}(x, y)$ : Merge the set containing  $x$  and the set containing  $y$ .
- $\text{FindSet}(x)$ : Returns a unique identifier for the set containing  $x$ .

1	2	3	4	5	6
$\{v_1\}$	$\{v_2\}$	$\{v_3\}$	$\{v_4\}$	$\{v_5\}$	$\{v_6\}$

$\text{FindSet}(v_5);$       Returns 5.

# The solution

We need to use a **union-find** data structure, also known as a **disjoint-set** or **merge-find** data structure. It supports the following operations:

- $\text{MakeUnionFind}(X)$ : Makes a new union-find data structure containing a 1-element set  $\{x\}$  for each element  $x \in X$ .
- $\text{Union}(x, y)$ : Merge the set containing  $x$  and the set containing  $y$ .
- $\text{FindSet}(x)$ : Returns a unique identifier for the set containing  $x$ .

<b>1</b> $\{v_1, v_2\}$	3 $\{v_3\}$	4 $\{v_4\}$	5 $\{v_5\}$	6 $\{v_6\}$
----------------------------	----------------	----------------	----------------	----------------

$\text{Union}(v_1, v_2);$

# The solution

We need to use a **union-find** data structure, also known as a **disjoint-set** or **merge-find** data structure. It supports the following operations:

- $\text{MakeUnionFind}(X)$ : Makes a new union-find data structure containing a 1-element set  $\{x\}$  for each element  $x \in X$ .
- $\text{Union}(x, y)$ : Merge the set containing  $x$  and the set containing  $y$ .
- $\text{FindSet}(x)$ : Returns a unique identifier for the set containing  $x$ .

1 $\{v_1, v_2\}$	4 $\{v_4\}$	7 $\{v_3, v_5\}$	6 $\{v_6\}$
---------------------	----------------	---------------------	----------------

$\text{Union}(v_3, v_5);$

# The solution

We need to use a **union-find** data structure, also known as a **disjoint-set** or **merge-find** data structure. It supports the following operations:

- $\text{MakeUnionFind}(X)$ : Makes a new union-find data structure containing a 1-element set  $\{x\}$  for each element  $x \in X$ .
- $\text{Union}(x, y)$ : Merge the set containing  $x$  and the set containing  $y$ .
- $\text{FindSet}(x)$ : Returns a unique identifier for the set containing  $x$ .



$\text{Union}(v_4, v_2);$

Note that  $\text{Union}$  may affect set identifiers unpredictably!

# The solution

We need to use a **union-find** data structure, also known as a **disjoint-set** or **merge-find** data structure. It supports the following operations:

- $\text{MakeUnionFind}(X)$ : Makes a new union-find data structure containing a 1-element set  $\{x\}$  for each element  $x \in X$ .
- $\text{Union}(x, y)$ : Merge the set containing  $x$  and the set containing  $y$ .
- $\text{FindSet}(x)$ : Returns a unique identifier for the set containing  $x$ .

42

$\{v_1, v_2, v_4\}$



$\{v_3, v_5, v_6\}$

$\text{Union}(v_5, v_6);$

Note that  $\text{Union}$  may affect set identifiers unpredictably!

# The solution

We need to use a **union-find** data structure, also known as a **disjoint-set** or **merge-find** data structure. It supports the following operations:

- $\text{MakeUnionFind}(X)$ : Makes a new union-find data structure containing a 1-element set  $\{x\}$  for each element  $x \in X$ .
- $\text{Union}(x, y)$ : Merge the set containing  $x$  and the set containing  $y$ .
- $\text{FindSet}(x)$ : Returns a unique identifier for the set containing  $x$ .



$\text{FindSet}(v_2)$ ;      Returns 42.

Note that Union may affect set identifiers unpredictably!

# The solution

We need to use a **union-find** data structure, also known as a **disjoint-set** or **merge-find** data structure. It supports the following operations:

- $\text{MakeUnionFind}(X)$ : Makes a new union-find data structure containing a 1-element set  $\{x\}$  for each element  $x \in X$ .
- $\text{Union}(x, y)$ : Merge the set containing  $x$  and the set containing  $y$ .
- $\text{FindSet}(x)$ : Returns a unique identifier for the set containing  $x$ .

42

$\{v_1, v_2, v_4\}$



$\{v_3, v_5, v_6\}$

$\text{FindSet}(v_5);$

Returns .

Note that Union may affect set identifiers unpredictably!

$\text{MakeUnionFind}$  takes  $O(|X|)$  time, and  $\text{Union}$  and  $\text{FindSet}$  take  $O(\log |X|)$  time. (It is also possible to add elements dynamically, but we won't need to.) So if we use this for  $\mathcal{C}$ ...

# Implementing Kruskal's algorithm: Third time lucky!

---

**Algorithm:** KRUSKAL

---

**Input** : Connected weighted graph  $G = ((V, E), w)$  in adjacency list form.

**Output** : A minimum spanning tree for  $G$ .

- 1 Sort the edges by weight as  $e_1, \dots, e_m$ , with  $w(e_1) \leq \dots \leq w(e_m)$ .
  - 2 Let  $T \leftarrow (V, \emptyset)$  be the empty tree on  $V$ .
  - 3 Let  $\mathcal{C} = \text{MakeUnionFind}(V)$ .
  - 4 **for**  $i = 1$  to  $m$  **do**
  - 5     Write  $e_i \rightarrow \{u_i, v_i\}$ .
  - 6     **if**  $\mathcal{C}.\text{FindSet}(u_i) \neq \mathcal{C}.\text{FindSet}(v_i)$  **then**
  - 7         Let  $T \leftarrow T + e_i$ .
  - 8         Call  $\mathcal{C}.\text{Union}(u_i, v_i)$ .
  - 9 Return  $T$ .
- 

Now line 3 takes  $O(|V|)$  time, and each iteration of lines 6 and 8 takes  $O(\log |V|)$  time.


So overall, since  $G$  is connected and  $|E| \geq |V| - 1$ , the running time is  $O(|E| \log |V|)$  — exactly what we got from Prim's algorithm!

## Non-examinable: Borůvka's algorithm

Neither Kruskal's algorithm and Prim's algorithm parallelise effectively.

But Borůvka's original algorithm, from 40 years earlier, works nicely.

At each step, it **simultaneously** finds and adds the cheapest edge out of **each component** of the output tree  $T$ .




## Non-examinable: Borůvka's algorithm

Neither Kruskal's algorithm and Prim's algorithm parallelise effectively.

But Borůvka's original algorithm, from 40 years earlier, works nicely.

At each step, it **simultaneously** finds and adds the cheapest edge out of **each component** of the output tree  $T$ .




## Non-examinable: Borůvka's algorithm

Neither Kruskal's algorithm and Prim's algorithm parallelise effectively.

But Borůvka's original algorithm, from 40 years earlier, works nicely.

At each step, it **simultaneously** finds and adds the cheapest edge out of **each component** of the output tree  $T$ .




## Non-examinable: Borůvka's algorithm

Neither Kruskal's algorithm and Prim's algorithm parallelise effectively.

But Borůvka's original algorithm, from 40 years earlier, works nicely.

At each step, it **simultaneously** finds and adds the cheapest edge out of **each component** of the output tree  $T$ .




## Non-examinable: Borůvka's algorithm

Neither Kruskal's algorithm and Prim's algorithm parallelise effectively.

But Borůvka's original algorithm, from 40 years earlier, works nicely.

At each step, it **simultaneously** finds and adds the cheapest edge out of **each component** of the output tree  $T$ .




## Non-examinable: Borůvka's algorithm

Neither Kruskal's algorithm and Prim's algorithm parallelise effectively.

But Borůvka's original algorithm, from 40 years earlier, works nicely.

At each step, it **simultaneously** finds and adds the cheapest edge out of **each component** of the output tree  $T$ .




## Non-examinable: Borůvka's algorithm

Neither Kruskal's algorithm and Prim's algorithm parallelise effectively.

But Borůvka's original algorithm, from 40 years earlier, works nicely.

At each step, it **simultaneously** finds and adds the cheapest edge out of **each component** of the output tree  $T$ .



Most modern algorithms for minimum spanning tree are variants of Borůvka's algorithm...and they use a union-find data structure to keep track of the components! So it is useful, after all.